

1 Customizing SQLSmith for SQL Query Generation

1.1 SQLSmith

SQLSmith is an open-source tool for randomized generation of SQL queries, designed originally for database fuzz testing.

SQLSmith is implemented as a grammar-driven random generator. It relies on three main components:

1. **grammar**: grammar production, defines the syntactic structure of queries (SELECT, DELETE, MERGE, FROM, JOIN, etc.);
2. **expr**: value expressions production, generates expression including values, predicates, comparisons, functions, and aggregations;
3. **relmodel**: models relations, columns, and aliases such as `ref_0`, `subq_0`, and `c0` which SQLSmith uses internally during generation.

1.2 Motivation for Customization

The goal is to adapt SQLSmith to generate a dataset of *realistic and structurally controlled* SQL queries, suitable for being used in ProvSQL and LLM fine-tuning. This required constraining or removing the following behaviours:

- automatically generated aliases (`ref_0`, `c12`, etc.);
- nonstandard system functions (distributed functions, probability functions etc.);
- TABLESAMPLE, LATERAL, and CTEs;
- nested aggregates (e.g., `SUM (MIN (x))`);

Three source files were modified: `grammar.cc`, `expr.cc`, `expr.hh`, `postgres.cc` and `relmodel.hh`.

1.3 Modifications to `grammar.cc`

The file `grammar.cc` controls the high-level structure of SQL queries. The following customizations were applied:

- disabling constructs that introduce uncontrolled complexity:
 - `table_subquery`;
 - `lateral_subquery`;
 - `table_sample`;
 - subqueries generated via `WITH` (CTEs).
- rewriting the `table_ref::factory` method to produce only real base tables or joins between base tables;
- simplifying the `SELECT` list by removing automatically generated derived column aliases (`c0`, `ref0`, ...);

- enforcing the use of only `SELECT` statements (disabling `DELETE`, `UPDATE`, `MERGE`, etc.);
- allowed only inner joins by modifying the `join_cond::factory` method to always set the join type to `inner`.
- adding support for `GROUP BY` clauses when aggregates are detected in the `SELECT` list. This was done by:
 - detect aggregates in `select_list` and mark the parent `query_spec` with a pending `GROUP BY` clause;
 - generate a `GROUP BY` clause in `query_spec::out` if marked as pending.
 - ensuring that only non-aggregate columns are included in the `GROUP BY` clause by modifying the relevant section in `query_spec::factory`.
 - avoid the use of `DISTINCT` when a `GROUP BY` clause is present.

1.4 Modifications to `expr.cc`

The file `expr.cc` defines how expressions, predicates, and function calls are generated. Some modifications were applied, as:

- **Disabling subqueries in expressions** (`atomic_subselect`);
- **Disabling window functions** (`window_function`);
- **Avoiding nested aggregate functions** by modifying the `funcall` constructor:

```

1
2 if (agg && dynamic_cast<funcall*>(p->pprod)) {
3     fail("nested aggregate not allowed");
4 }

```

- **Removing complex or system-defined functions** from the generative process.
- **Removing exists predicates** from the generative process because it is not supported in ProVSQL.
- **Modifying the set of available operators** to use only basic operators in comparisons.

```

1         bool name_ok =
2         name == "=" ||
3         name == "<" ||
4         name == ">" ||
5         name == "<=" ||
6         name == ">=";

```

- Modified function `sympole_join` to allow joins only on ids, refs, or integer columns.

1.5 Modifications to `expr.hh`

The file `expr.hh` defines various expression types. The following change was applied:

- **Modifying the `is null` predicate** to always generate `IS NOT NULL` predicates. This was a problem because ProvSQL does not give any provenance information for `IS NULL` predicates.

```

1     virtual void out(std::ostream &out) {
2         // modified to avoid IS NULL
3         //out << *expr << " is " << negate << "NULL";
4         out << *expr << " is not NULL";
5     }

```

1.6 Modifications to `postgres.cc`

The file `postgres.cc` defines PostgreSQL-specific functions and types. The following changes were applied:

- removing PostgreSQL aggregates and routines from the generative process, such as `pg_catalog.random()`, `pg_catalog.setseed()`.

1.7 Modifications to `relmodel.hh`

The file `relmodel.hh` defines the representation of tables, relations, and aliasing. SQLSmith generates internal aliases such as `ref_0`, `ref_1`, and derived column names such as `c0`, `c1`.

The following changes were applied:

- aliases in the table naming and routine naming are disabled, using the table name directly.

This transforms outputs like:

```

1 FROM public.partsupp
2 FROM system.partsupp
3 SELECT pg_catalog.min()

```

into:

```

1 FROM partsupp
2 FROM partsupp
3 SELECT MIN()

```

1.8 Resulting Behaviour

After applying these modifications, SQLSmith produces:

- only `SELECT` queries (no CTEs, no `MERGE/UPDATE/DELETE`);
- joins between real base tables only;
- simple and readable `SELECT` lists without synthetic aliases;
- aggregate functions without nesting;